

Exercise task:

*Hello, I am your new client and I need your help with my issue. I have thousands and thousands of papers in the boxes back in my garage. The important part is that there is always some number on each paper. Now I would like to somehow extract all these numbers and for example store them somewhere, in a text file, Excel, it doesn't really matter. Can you help me with that?*

*Would it be possible to provide and show some solution without even sending you any of my data?*

*What would you need to know for obtaining a complete scenario of the use-case for some "robust" solution?*

Exercise - Computer Vision Scientist

Problem: Text Extraction - „I have a number printed on a paper and I want to read it.“

Goal:

- take some data, try some approaches, produce some code,
- get some results - then come and show us your solution and have a chat around it - show how you think about a specific problem, how you are able to explain what approach you used and why, think about the limitations of the approach and how things could be improved if there was more time, what you think of the results and if they make sense, etc.

Philosophy:

- the actual result and code are not that important – the journey there and potential future paths are more important
- if you are able to follow-through with some ideas then great, if you just start something and have a clear idea on how to proceed that is also useful
- the expectation is that you will spend an evening or two with the task (but there are no bounds to proactivity if you enjoy playing with the problem)

Data:

- feel free to grab any relevant data set

Tools:

- use whatever you are comfortable in (R, Python, Matlab, Java, SQL,...) or feel free to use it as an opportunity to try out a new language
- it is not a contest in finding the best black-box library and blindly using it
- own solutions are preferred even if they are simple
- how to present: up to you – slides, walking through code, drawing, ...